

PRODUCT MANAGER TAKE-HOME · NOMAD VENTURE STUDIO / JOBESCAPE

Jobescape: Audience, Competitors, the Challenge & the Economics

A research dossier — the evidence, the thesis, and the plan behind
Parts A–D.

Prepared by Aliaskar Bekishev · 2026-07-28 · Built with a team of AI research agents (on-brand: Jobescape teaches working *with* AI)

Three-step method: broad research (scientific · social · market) → synthesis → deep-dive. Every external claim carries a source; internal-data gaps are flagged, not faked.

Contents

Executive summary

- [The thesis](#)
- [What the evidence establishes](#)
- [The numbers](#)
- [What is measured, and what is inferred](#)

How this was done

- [Three passes](#)
- [Two things built directly, not researched](#)
- [How the numbers are kept honest](#)
- [Evidence standards](#)

Part A — Audience & Product

- [A.0 — The funnel, decoded \(primary evidence\)](#)
- [A.1 — Audience segments](#)
- [A.2 — Does the product meet the need? \(expectation vs. reality\)](#)

Part B — Competitors

- [B.1 — The competitive map](#)
- [B.2 — Teardown highlights \(what each does well, and the one thing worth taking\)](#)
- [B.3 — What to take, in priority order \(and why this order\)](#)

Part C — The "Challenge" release

- [C.1 — Release analytics](#)
- [C.2 — Verdict: No, the release was not a success](#)
- [C.3 — What's next: fix the placement, then the mechanic](#)
 - [The original phasing, still valid for the mechanic itself](#)
- [C.4 — Prototype \(Task 4\)](#)

Part D — Subscription economics

- [D.1 — Total Net LTV \(one-year horizon\)](#)
- [D.2 — A/B plan-upgrade break-even](#)

The economics thesis (why retention is the moat)

- [Funnel reconciliation — 1,000 signups/day → annual net revenue](#)
- [The tail is the business — and AI apps churn through it faster](#)
- [Retention beats the funnel \(and it's safer\)](#)
- [The asymmetric tail risk — and the one move that fixes three problems](#)

Appendix

- [A — Jobscape quiz funnel, full decode \(quiz_version v7.0.8\)](#)

Executive summary

The thesis

Jobscape has built an excellent machine for selling a product it hasn't finished building.

Acquisition works. Activation was never built. And the economics depend almost entirely on the retention tail that activation would protect.

The measured fact underneath it: **50.2% of paying subscribers never complete a single lesson** — not in the Challenge, anywhere in the product. Everything below is a facet of that one seam.

What the evidence establishes

The funnel sells fear, and it sells well. The 37-step quiz profiles the user, then systematically pokes pain — *"What scares you most about AI and your career?"* — before stacking FOMO teasers and presenting payment at the peak of it. That is rational positioning, not cynicism: 52% of US workers are more worried than hopeful about AI (Pew), and 66% of leaders say they won't hire someone without AI skills (Microsoft/LinkedIn).

But fear only closes the sale. Habit and relevance are what retain, and self-paced learning is a churn minefield — MOOC completion runs 5–15%, and the drop-off is motivational rather than cognitive. A product built on this funnel has to convert fear into habit quickly, or it inherits that churn.

The buyer is a decade older than the internal read. Around 60% of paying subscribers are 45+, and the largest single group is men aged 55+ — 21.3% of all buyers. Churn falls steadily with age: 31.1% at 18–24 against 9.5% at 45–54. The oldest buyers are the best buyers. That makes streak-and-badge mechanics a design aimed at the least valuable 9% of the book.

The Challenge was the right instinct, wrongly built. Measured across 9,956 paying subscribers: 89% never start it. Among those who do, the target metric doesn't move. Takers retain at 60.0% on D1 against **59.6% for people who reached the same screen and walked away** — a gap of +0.5 points, $p = 0.82$.

The reason is mechanical. The hypothesis was "one day = one skill builds a daily habit," but nothing gates the content to one day, so **27% of finishers completed the whole "7-day" challenge in a single sitting**. Version one didn't fail its hypothesis. It never tested it.

The content is free elsewhere, and the funnel is already cloned. Structural twins run the identical playbook — Coursiv, and Iro AI which is a literal "Duolingo for AI" — while ChatGPT, YouTube and newsletters give the material away. Defensibility has to come from scaffolding, feedback and accountability: the AI tutor as a practice surface, which is the one thing free substitutes cannot provide.

The acute risk is billing, not content. Ratings are strong — 4.7 on the App Store, 4.5 across ~8,000 Trustpilot reviews — yet complaints cluster on surprise charges and refunds gated on course completion. In this cohort, **11.3% of all cancellations are payment failures, chargebacks or disputes.** The FTC sued a peer quiz-funnel studio, Genesis, over this exact model in June 2026. The high rating hides the exposure.

One correction worth flagging, because it cuts against the reviews: cancellation itself is **self-serve**. The event log records 1,358 uses of the Manage Subscription flow. The trap is the refund, not the cancel button.

The numbers

	1-WEEK	4-WEEK	12-WEEK	BLENDED
Net LTV, one-year horizon, 12% fee	\$60.42	\$123.70	\$162.15	\$125.06

The A/B plan-upgrade test breaks even at a **4.9% take-rate** — 8.3% under the conservative horizon. Both are low bars, and the only thing at risk is the \$3.60 second upsell it replaces. Recommend running it.

What is measured, and what is inferred

Measured from source data. All of Part C, from BigQuery across 9,956 paying subscribers. Part A's segments, from k-modes clustering over 38,071 quiz respondents. Part D, from a model built over the supplied plan table and cross-checked against the observed plan mix. Every published figure in Parts C and D is re-derived from the raw exports by a verification script — 91 assertions, all passing.

Inferred from public evidence. Part B's competitor teardowns, and the expectation-versus-reality gap in Part A — sourced from public reviews plus my own walkthrough as a paying customer.

What this data cannot settle. Whether the Challenge was causal. There is no hold-out group and no pre/post period anywhere in the warehouse; the release shipped to 100%. The causal reading here is an observational reconstruction, and it only worked because a large "looked at it and left" group happened to exist as a natural control. That is luck, not design — which is why fixing it is the *first* recommendation in Part C rather than the last.

How this was done

The brief grades quality of thinking, and it is on-brand for a company teaching "use AI at work" to see AI used well. So this ran as a small research operation rather than a single sitting — and, more importantly, with a rule that every number had to survive being checked.

Three passes

Broad research. Five parallel agents swept the field in three directions: the science of habit formation and adult learning; the social picture of AI-at-work anxiety and reskilling; and the market — AI-upskilling edtech, the quiz-funnel model, and competitors direct and indirect. Each returned a sourced report with confidence flags.

Synthesis. The reports were cross-read to find the one thesis connecting all four parts, and to pick which threads justified going deeper.

Deep dives. Four focused agents: two competitor teardown sets, a Challenge v2 and prototype design, and a funnel economics model.

Two things built directly, not researched

The quiz funnel was decoded end to end. Rather than clicking through 37 screens, the funnel's public constructor API was queried to reconstruct every page, question, answer option, the variable each answer writes, and the single branch point. That is primary evidence for Part A, not observation.

The subscription economics were modelled from scratch over the three plans, and cross-checked against the plan mix actually observed in the data.

How the numbers are kept honest

This is the part that matters most, so it is stated plainly.

Part C runs on one BigQuery export. Five queries produced the data;

`13_comparison_groups.sql` returns one row per paying subscriber, takers and non-takers alike — 9,956 rows. Three earlier queries were written against an assumed schema, superseded once the real schema was probed, and are shipped as-is with their `△ADJUST` markers, because they are the honest record of the working.

The analysis is six Python scripts, pure standard library. One settles the questions that had to be answered *before* the engagement tiers were chosen — where the retention curve breaks, how much observation window exists, whether challenge 338 really is "the Challenge". The others produce the release analysis, the supplementary tests, and the cohort baselines.

Every published figure is re-derived and asserted. `05_qa.py` recomputes 91 numbers straight from the raw CSVs and fails loudly if any documented figure has drifted. It exists because these numbers had to survive being copied into four places — analysis notes, this dossier, the deck, the submission — and hand-copying is where quiet errors live. It has already earned its keep: it caught a 59.557% that had been rounded to "59.5%" and propagated to eight files.

Part A's segments were derived, not asserted. The first version was a hand-written rule chain — six segments authored from a jobs-to-be-done framework, then counted. It was coherent and unfalsifiable: the structure came from the analyst, so no data could contradict it. It was replaced with k-modes clustering over 38,071 respondents. The clustering rejected two of the six, and surfaced one real segment the framework had missed entirely. That result is reported in full in Part A, including the part that is inconvenient.

Evidence standards

Every external claim carries a source, and figures from secondary aggregators or vendor blogs are flagged as softer than primary data. Assumptions are stated rather than buried, particularly in the economic model. Where a claim rests on internal data, the script that computes it is named. Where the data cannot answer a question — most importantly whether the Challenge release was causal, which no observational dataset settles without a hold-out — that is said plainly rather than papered over.

Part A — Audience & Product

A.0 — The funnel, decoded (primary evidence)

Jobscape's quiz was reconstructed end-to-end from its public constructor API — all **37 pages, every question, every answer option, and the one branch point**. It is a textbook "diagnose the pain, then sell the cure" funnel in four movements:

1. **Hook + segment on prior AI use.** Page 1 asks *"Have you ever used Claude?"* and branches on `used_ai` — the funnel's only fork. "Yes" → a FOMO card ("You're ahead — but this is just the start"); "No" → a reassurance card ("You're right on time"). Same destination, tailored emotion.
2. **"My Profile" — biometrics & intent.** why-Claude · work status · age · gender · goal → a *"Did we get everything right?"* summary. This is the raw material for demographic segmentation.
3. **"Challenges" — poke the pain, then teaser-sell.** Three pain blocks, each followed by a persuasion teaser: career fear (*"What scares you most about AI and your career?"*) → *"There is nothing to worry about"* + the Harvard "humans with AI will replace humans without AI" quote; document pain (*"used AI then*

spent just as long fixing it?") → "Create polished content"; build-an-app pain → "Building an app has never been this easy."

4. **"Personalization"** → **commitment** → **paywall**. Learning-style, time-commitment, certification desire, *"what stopped you before?"*, an AI-profile reveal, a fake-progress loader, **email + name capture**, a "Your Personal Plan" teaser, then the **selling page** and a second "chase" offer.

The read: the funnel's *product* is emotional — it manufactures urgency (FOMO), self-identified pain, sunk-cost investment (35 taps of personalization), and then presents payment at the peak of aroused anxiety. Full page-by-page map in the Appendix.

A.1 — Audience segments

How these were produced. The first version of this section was a framework: six segments authored from a jobs-to-be-done model, then counted by a hand-written rule chain. It read well and it could not be wrong — the structure came from me, so no answer in the data could contradict it.

So it was replaced with a real segmentation. k-modes clustering over **38,071 quiz respondents × 10 categorical answers**, from a 49,528-respondent export. k-modes rather than k-means because the answers are nominal: the average of "35–44" and "business owner" is not a quantity, and one-hot plus Euclidean distance quietly assumes every pair of answers is equally far apart — which is the assumption under test.

The honest headline is that there is no elbow. Within-cluster cost falls smoothly across $k = 2 \dots 8$ (5.09, 4.79, 4.67, 4.49, 4.40, 4.31, 4.20) with no knee. This audience is a continuum, not a set of natural types. Six is a resolution chosen and defended, not a count the data forced — and any segmentation of this audience should be presented that way.

SEGMENT	WHO THEY ARE (CLUSTER SIGNAL)	SHARE	STRONGEST LIFT
The Striver	senses a gap and is closing it deliberately — growth-motivated, 45–54, male-skewed	28.4%	"somewhat ready, could be behind" 1.9×
The Latecomer	the oldest buyers; hasn't started, watching the field move without them	22.0%	"falling behind as others move faster" 2.2× · never used Claude 1.9×
The Adept	already gets real value from AI and wants more leverage — female-skewed	20.9%	"AI already helps me a lot" 2.5×
The Optimiser	mid-career professional using Claude on real work, wants speed	12.3%	aged 35–44 2.6×
The Founder	learning AI to build something of their own	9.0%	"start my own business" 3.1×
The Switcher	between jobs, changing direction, starting from zero	7.4%	"exploring options" 4.4×

Derived by *k*-modes clustering over 38,071 quiz respondents × 10 categorical answers (`part-a-audience/analysis/cluster_quiz.py`), not authored. Lift = share inside the cluster ÷ share in the population. There is no elbow — within-cluster cost falls smoothly across $k = 2 \dots 8$ — so this audience is a continuum and six is a defended resolution, not a count the data forced.

Why the clusters split on motive rather than age: occupation-level AI adoption ranges ~12–65%, and self-assessed *confidence* — not age — is the strongest predictor of use (HBS). The data agrees: the sharpest single signal is **work status** ("exploring options", 4.4× lift), not a demographic band.

The buyer-side cut comes from `subscribe_events` (9,956 paying subscribers), read against the quiz clusters. An earlier hand-written rule chain produced a different six: clustering rejected two of them and surfaced one — The Adept — that the framework had missed entirely.

Three findings that overturned my priors:

1. **The audience is old, not young.** ~62% of quiz-takers and ~60% of buyers are 45+; the largest single cell among buyers is **men aged 55+ (21.3%)**. Buyers are 59.7% male. The AI-adoption age skew was the wrong prior for *this* funnel.
2. **Older buyers are the better buyers.** Two-week unsubscribe falls monotonically with age: 31.1% (18–24) → 19.7% → 14.5% → 9.5% (45–54) → 9.7% (55+). The loyal core is the oldest half of the book — which makes gamified streak mechanics a design aimed at the 9% that churns hardest.
3. **The goal a buyer arrives with predicts whether they stay.** "Work faster" (30.3% of buyers) and "feel more confident with AI" (20.0%) churn at ~11%; **"get a quick side hustle" churns at 37.2%**, "gain flexibility / work remotely" at 27.4%, "unlock better income" at 25.4%. The income-and-freedom intents are ~12% of buyers and churn at 2.5–3.5×. **Every creative that promises a side hustle buys a customer who leaves** — the positioning inconsistency of A.2, priced.

Also notable: only 47% of quiz-takers name a career fear at all — the rest call AI an opportunity. The fear-heavy funnel over-indexes on one half of its own audience.

A.2 — Does the product meet the need? (expectation vs. reality)

The *reality* column combines **public user evidence** — App Store 4.7 (~1.5K), Trustpilot 4.5 (~8K), plus the Coursiv twin as proxy (`research/deep-05`) — with my own walkthrough as a paying customer (I bought the 4-week plan for \$15.19 and screenshotted the funnel, paywall, product and cancel flow; notes in `part-a-audience/materials/walkthrough/observations.md`). **Tell:** ratings are high *and* billing complaints are constant — an aggressive review-solicitation funnel over a dispute-

heavy monetization model; the score is a marketing asset, not proof the gap is closed.

Three things I could only confirm by paying for it:

- The paywall's **"personal AI mentors"** is a single bot with a stock-photo face and a human name ("Mark"), present on every page as an "Ask Mark" button.
- The paywall's **"24/7 support chat"** is a **metered 5-credit** AI chat.
- **Live pricing differs from the supplied plan table** — a "61% intro offer" with a 9:46 countdown: \$6.93 / \$15.19 / \$25.99 intro, renewing at \$38.95 / \$38.95 / \$66.65. Part D uses the brief's figures as instructed; the real dynamic pricing is flagged as an observation.

And one correction to the review-sourced picture: the widely-repeated *"bot-only cancellation"* claim is **wrong for cancellation**. The paywall fine print names a self-serve **Manage Subscription** tab, it exists, and the event log records **1,358 uses of it**. The trap is the **refund** (gated on finishing the plan in ~31 days), not the cancel. Worth being precise about — it's the difference between a dark pattern and a bad guarantee.

LENS	ENTRY-POINT PROMISE	WHAT USERS SAY THEY GOT (REVIEWS)
Expected	~\$25 "trial"; personalized plan; expert instruction; exclusive tools	A bigger bill (+\$88 "AI support" in fine print , "\$100 not a trial," silent renewals, a bank-flagged-as-fraud charge); content "you can just ask ChatGPT," repackaged free tools, anonymous team
Doubts	"Will this work for me? Am I behind?"	Often <i>reinforced</i> — "shallow vs. free alternatives"
Promised, not delivered	"Earn or money-back" guarantee; a freelance <i>income</i> ("land clients in 3 months")	Refund gated on finishing 14 modules / the plan in ~31 days ; outcome - buyers feel misled ("gurus promising 10k/mo")

Where the main gap is, and the risk. I came in expecting the core gap to be the *motivation switch* (fear-sale → slow habit product) — and that's real. But the evidence relocates the *acute* risk one layer down: the **price-and-exit gap** — users are billed more than they agreed, then can't easily get their money back (the refund weaponized as a completion trap — cancellation itself is self-serve). That's not just churn — it's **refunds, chargebacks, 1-star reviews, and regulatory exposure**: this is the exact model the **FTC sued Genesis (a peer) over in June 2026** (ARL tail ≈ 3× annual net revenue, Part D). The 4.5–4.7 ratings **hide** it. Compounding it: a **positioning inconsistency** (jobscape.me "AI at work" vs. app "freelancing income" vs. a rotating App-Store title) pre-loads mismatched expectations. The fix is the same one Part D calls an economic moat and Part B ranks first — **transparent pricing + a guarantee that isn't a trap**.

Part B — Competitors

B.1 — The competitive map

How I split direct vs. indirect: *direct* = solves the same job (get better at using AI for work) for the same audience (non-technical professionals) in a substitutable format; *indirect* = competes for the same attention, budget, or acquisition channel while differing in subject or format. Substitutes (free tools) are a third, and — as the research shows — the *most dangerous* category.

COMPETITOR	TYPE	WHY IT COMPETES
Coursiv	Direct — structural twin	Same quiz → app → subscription funnel, same "learn AI for your job," same price band. The closest rival.
Iro AI	Direct — product twin	Literal "Duolingo for AI." Closest <i>product</i> , but weak commercially (soft funnel, ~\$60/yr, small brand).
Section (SectionAI)	Direct — premium incumbent	Same audience (AI for knowledge workers); credibility via Scott Galloway; now pivoting to an AI-coach SaaS.
Outskill	Direct — high-ticket funnel	AI upskilling via free workshop → \$1–3k cohort; bids in the same ad auctions.
Be10x	Direct (near-future)	Funnel-mechanics twin (₹9 tripwire → mastery upsell); India-first today, US launch Nov 2025 .
The Rundown / Rundown University	Direct — media flywheel	Teaches applied AI; acquires via a 2M+ owned newsletter, not the ad auction.
Duolingo	Indirect — the template	Different subject, but <i>the</i> engagement engine Jobscape imitates (streaks/daily habit).
Headway	Indirect — GTM twin	Same Meta-ads → quiz → paywall → gamified playbook, applied to microlearning/books.
Mimo	Indirect — UX exemplar	Gamified coding app; best-in-class onboarding & paywall mechanics to copy.
Sololearn	Indirect → near-direct	Gamified mobile app now explicitly teaching AI skills (prompt engineering, ML).
ChatGPT / Claude themselves	Indirect substitute	<i>The</i> free self-teaching path — learning by doing, at zero cost.
YouTube · Coursera · Udeemy	Indirect substitute	Free/cheap credible content covering the same skills.
AI newsletters (Superhuman, Mindstream)	Indirect	Own "apply AI in 5 minutes" — almost Jobscape's verbatim promise.
Employer L&D	Indirect substitute + expansion path	Free-to-user corporate AI training; also Jobscape's future B2B tier.

B.2 — Teardown highlights (what each does well, and the one thing worth taking)

Structural twins & ad-funnels. *Coursiv* runs the identical machine and has branded its retention arc as a **finite "28-Day AI Challenge" with a completion certificate** — turning vague AI dread into a *closable* loop with a natural renewal moment; its **275 BBB complaints over undisclosed recurring pricing** are the avoidable failure mode. *Iro AI* ships a "Prompt Lab" (graded practice against a live model) that kills the "too passive/basic" complaint. *Outskill* proves the *same fearful buyer will pay \$500–3,000* for a live time-boxed event. *Be10x* uses a **~\$1 commitment tripwire** to manufacture sunk-cost show-up commitment from low-intent ad traffic.

Premium & media incumbents. *Section* has pivoted from courses to an AI-coach SaaS — the loudest signal in the field: even the premium incumbent admits courses alone don't retain, so it's bolting on a recurring *tool* with saveable, reusable use-cases. *The Rundown* grows a 2M+ newsletter via **beehiiv Boosts at ~\$2/subscriber** (moved ~50% of budget off Meta) and prices University at the **"expense-it" \$1,008 threshold** — **~70% bill it to employer L&D**. *Mimo* is the UX playbook: a ~15-step commitment-escalating onboarding and the **"Blinkist paywall" experiment (trial-timeline screen + pre-expiry reminder) that doubled trial opt-ins (+100%) and cut 24h cancellations 25%**.

The template & the substitutes. *Duolingo* proves the streak is the moat (7-day streak = 2.4–3.6× continuation). But the sharpest strategic fact in Part B is that the content itself is free: ChatGPT/Claude, YouTube, and newsletters give away exactly what Jobscape sells. Jobscape cannot win on content — only on scaffolding, feedback, and accountability (the AI tutor as a *practice surface*, the one thing free substitutes structurally can't provide).

Reconciling this with Part C (worth stating, because it looks like a contradiction). Part B says the streak is Duolingo's moat; Part C recommends *removing* gamification from the Challenge. Both are right, because Duolingo's real lesson is **daily pacing**, not the flame. What retains is that today's lesson is available today and tomorrow's is not — Jobscape's Challenge copied the pageantry (a streak counter, a badge, a certificate) and skipped the pacing, which is why its streak displays 0 and 27% of finishers binge the whole thing in a day. And Duolingo's audience is not Jobscape's: ~60% of Jobscape's buyers are 45+, and the under-25s who respond best to streak mechanics are 9% of the book and churn at 31%. Take the rhythm, leave the confetti.

B.3 — What to take, in priority order (and why this order)

The order follows the money: de-risk the fragile economics first, then build retention, then defensibility, then diversify acquisition, then expand monetization. (Ties directly to Part D: the value is a fragile recurring tail, so protecting it beats squeezing the funnel.)

1. **Copy Mimo's "Blinkist paywall" verbatim** (trial-timeline + pre-expiry reminder). *Why first:* proven **+100% trial opt-in / -25% early cancels**, near-zero build cost, and it directly attacks the refund/chargeback risk that threatens the recurring tail. Highest ROI, lowest effort.
2. **Adopt Coursiv's branded finite "Challenge" as the retention container — with honest recurring-price disclosure.** *Why second:* it reframes the whole product around the habit loop (pairs with the Part C v2) *and* defuses the FTC/Genesis-style existential risk. Retention + compliance moat in one move.
3. **Build the practice surface: Iro's "Prompt Lab" (graded practice) + Section's role-specific co-pilot with saveable use-cases.** *Why third:* this is the moat against the real threat — *free substitutes*. It converts the existing chatbot from commoditized content delivery into feedback + switching cost. Medium effort, core to defensibility.
4. **Launch an owned newsletter / daily in-app brief as a CAC hedge (à la The Rundown) + frame the annual plan as "expensable."** *Why fourth:* it reduces 100%-paid-ads dependence and improves Part D break-even the moment blended CAC drops — but it compounds slowly, so start early, harvest later.
5. **Test the offer ladder: Be10x's ~\$1 commitment tripwire (lift activation) + Outskill's high-ticket live "AI Mastermind" (lift ARPU).** *Why last:* monetization expansions worth testing *after* retention and differentiation are fixed — a lower entry to raise activation, and a higher-LTV offer for paid ads to bid against.

Part C — The "Challenge" release

The Challenge delivers the course as a **7-day, one-skill-a-day** arc. Target metric: **D1 retention**, with unsubscribe, completion, D3, D7 and CSAT secondary.

Data: BigQuery `persona-496908.sql_assessment`, pulled 2026-07-27. **9,956 paying subscribers** whose first app activity falls in 2026-06-12 → 2026-06-25. The feature is `course_id = 338` (*7-Day Claude Challenge*, 6 languages), which accounts for **99.7% of all challenge starts**. Every figure below is reproduced by the scripts in `part-c-release-verdict/analysis/`.

C.1 — Release analytics

Engagement tiers, defined explicitly. High = started and completed **≥3 of the 8** challenge lessons. Low = started, **≤2** lessons. Didn't take it = never started. The cut sits at 3 because that is where the D7 curve steps (18% → 39%), and sweeping every possible threshold shows the retention gap growing monotonically while the unsubscribe gap stays at zero — so the conclusion does not depend on where the line is drawn.

The exposure ladder — the finding that reframes everything else:

STEP	USERS	% OF ALL PAYERS
Paying subscribers	9,956	100%
... saw the Challenge popup	3,140	31.5%
... viewed a Challenge page	2,109	21.2%
... started Challenge 338	1,111	11.2%
... completed ≥1 lesson	612	6.1%
... completed all 8	48	0.5%

89% of paying subscribers never start it; 67.5% never see a challenge surface at all. Once someone does see the popup, 94.6% click it — desire is not the constraint, distribution is.

Inside the feature: 45% of starters complete zero lessons. Half of those who finish lesson 1 never open lesson 2. Past lesson 3 the drop-off almost stops (survival climbs 51% → 90%).

The headline comparison (all groups anchored on their first app day, 95% Wilson intervals):

METRIC	HIGH (N=158)	LOW (N=953)	DIDN'T TAKE (N=8,845)
D1 (target)	69.6%	58.4%	21.7%
D3	57.0%	34.4%	10.0%
D7	39.1%	18.0%	6.9%
Unsubscribe	11.4%	11.4%	14.3%
CSAT	3.33	3.53	3.74

Read literally this says the Challenge triples D1. It does not.

C.2 — Verdict: No, the release was not a success

Four tests, and the feature fails all four.

Exposure-matched control (decisive). "Didn't take it" is mostly people never offered the feature. The fair control is people who reached the Challenge and declined it:

GROUP	N	D1
Started 338	1,111	60.0%
Joined 338, never started	230	60.4%
Viewed 338, never joined	764	59.3%

Takers vs exposed-but-didn't-start: +0.5 pts, $p = 0.82$. Holding total product activity constant, the effect stays at zero in four of five strata. Unsubscribe: -2.3 pts, $p = 0.12$, not significant.

Immortal-time bias. Only 24% of takers start on their first app day; the median starts on day +1 and 28% on day +3 or later. To start on day k you must already have survived to day k . Correcting for this alone cuts the apparent gap from +38 pts to +17.

A non-tautological dose-response. Among takers whose entire challenge activity happened on one day, a bigger day-0 dose predicts **worse** next-day retention (-9.2 pts, $p = 0.048$). Not flat — negative.

The mechanic ships, but only as a suggestion. Trying to open tomorrow's lesson today raises a dismissible warning — real day-logic, with no enforcement behind it and no event logging whether it was obeyed or clicked through. The hypothesis needs daily gating; a warning you can click past is not gating. **27% of the people who completed the "7-day" challenge did it in a single day; 46% within two.** Mean 1.46 lessons per active challenge day. Only 36.8% of starters who completed a lesson ever returned for a second challenge day.

v1 did not fail the hypothesis — it never tested it. What shipped was a short course with "7-Day" in the title.

The number that would have fooled the team. On the brief's own metric definition — D1 measured from the challenge start day — the Challenge posts **39.8%** against a product baseline of **26.0%**. Any dashboard built to spec reports a +14-point win. It is an artifact of two different clocks and two different populations. And there is no hold-out anywhere in the data: the release went to 100% with no control arm, so everything above is an observational reconstruction that only worked because a large "looked and left" group happened to exist.

A red flag that changes the recommendation: satisfaction falls as engagement rises — 3.74 (never took it) → 3.53 → 3.33 → **3.10 (finished it)**. Challenge content rates 3.25 while the same users rate the rest of the product 3.45.

C.3 — What's next: fix the placement, then the mechanic

The reframe. 94.6% of subscribers shown the Challenge popup clicked it; 68.5% were never shown it. Appeal is not the constraint — placement is. Making the Challenge the default post-onboarding path, shipped to 80% with a 20% hold-out, fixes the largest loss in the funnel and creates the control group v1 never had. Guardrail: Challenge CSAT is 3.25 against a 3.45 baseline, so watch refunds in the test arm.

And day 8 is a cliff. The Challenge and the personal plan are independent tracks — the Challenge gated and daily, the plan open and self-paced — and nothing connects them. A challenge day ends with tomorrow locked and nowhere to go. Route the user to the plan instead; 75.7% already start one, and nothing sends them there.

The original phasing, still valid for the mechanic itself

The reframe. The Challenge is treated as a retention feature. It should be the product's **activation path** — because **50.2% of paying subscribers never complete a single lesson anywhere**, and a finite, sequenced, one-thing-at-a-time path is exactly the right shape of on-ramp for a 45+ professional.

The counter-intuitive call: reach is the LAST fix, not the first. The obvious response to 11% reach is to promote it. But satisfaction falls with engagement, so scaling reach on 3.25-star content scales refunds — already 11.3% of unsubscribes are payment failures, chargebacks or disputes. Order: **instrument** → **mechanic** → **first ten minutes** → **quality** → **reach**.

PHASE	THE WORK
0 · Instrument	10–20% hold-out; redefine D1 as exposure-matched on a common clock; log day-unlock events; extend the window past first renewal
1 · Mechanic	Gate the content to one day at a time — this <i>is</i> the experiment. Plus a scheduled trigger at a user-chosen time, tomorrow named as a concrete task, and catch-up rather than punishment
2 · First 10 min	Find out what happens to the 45% who press Start and complete nothing (a bug hunt before a design task); guarantee a win inside 10 minutes on a pre-loaded example; bridge lesson 1 → 2
3 · Quality	Mine the 3,601 exit-reason and 1,295 bug events already in the warehouse; re-aim framing at the real buyer — competence, not confetti
4 · Reach	Make it the default post-onboarding path — gated on CSAT reaching the 3.45 product baseline

Remove: the "7-Day" name on an 8-lesson course; the certificate (23 downloads, 0.2%); any further gamification — the segment that responds to it is 9% of buyers and churns at 31%.

Pre-registered bar: exposure-matched D1 lift $\geq +5$ pts vs hold-out; starters completing ≥ 1 lesson 55% → 75%; guardrails on CSAT, refunds and total completions. Kill rule: if the daily gate ships with a proper hold-out and still shows no lift, kill the Challenge — the conclusion would be that this audience doesn't want a challenge, and activation needs a different shape.

C.4 — Prototype (Task 4)

Live: https://alyasska.github.io/Nomad_Venture_Studio_TA_C4/

A working React application built in Jobescape's own design system, so it reads as a feature of the product rather than a mockup of one. Five steps, each with its own

URL and each annotated on the page with the finding that produced it: **Commit** (pick the daily time) → **Today** (one day open, tomorrow locked with a live countdown) → **The win** (pre-loaded real task → drafted result → save it) → **Done** (a kept result; tomorrow named as a work task) → **Tomorrow** (the reminder fires, the gate opens). State persists, so the gate survives a refresh — a gate that resets is not a gate.

It borrows Jobscape's **visual language** — the blue, the rounded cards, the staggered lesson path — but not its **streak mechanic**. The buyer is 59.7% male and ~60% aged 45+, and the under-25s who respond best to streaks unsubscribe at 31.1%; the slot where the product's streak pill sits now counts the days the product has opened for you. The reward is competence, not celebration. The visual argument is part of the product argument.

How it improves the next release: it turns the plan into something you can put in front of five 50-year-old professionals on a Tuesday and watch. It de-risks the three decisions that would otherwise be argued in a room — whether a locked day reads as helpful or as being shut out of something you paid for (the one real risk in Phase 1), whether the pre-loaded first task produces a genuine "that's useful" inside ten minutes, and whether competence framing beats streak framing for this audience. Each costs a day to test on the prototype and a sprint to discover in production.

Part D — Subscription economics

Fully modeled from the raw plan data; reproducible via `part-d-economics/model/ltv_model.py`.

D.1 — Total Net LTV (one-year horizon)

Method. Per plan, model one subscriber: everyone pays the intro (survival $S_0 = 1$); survival to recurring period k is the running product of the given conversion rates $S_k = c_{01} \cdot c_{12} \cdot \dots$. Revenue = intro + recurring $\cdot \sum S_k$ + expected upsells (conv $\times$ price, one-time). Net = gross $\times$ (1 – 12% processor fee).

PLAN	MIX	GROSS LTV	NET LTV
1-WEEK	10%	\$68.66	\$60.42
4-WEEK	70%	\$140.57	\$123.70
12-WEEK	20%	\$184.26	\$162.15
Blended (mix-weighted)	100%	—	\$125.06

Worked example (4-week, the 70% plan): $\sum S_k = 2.400 \rightarrow$ recurring $39.99 \times 2.400 =$ \$95.98; + intro \$19.99; + upsells $(0.30 \times 69.99 + 0.12 \times 29.99) =$ \$24.60 $\rightarrow$ gross **\$140.57**; $\times 0.88 \rightarrow$ **\$123.70 net**.

Key assumptions (stated): 12% fee on *all* revenue incl. upsells; the 12 supplied Cmn transitions define the horizon (for the 4-week plan that is exactly 52 weeks — the anchor for "one year"); upsells fire once at checkout.

Sensitivity. The only wrinkle is the 12-week plan: billed 12 times it spans ~2.8 years, not one. Under a strict "payments charged within 52 weeks" reading it gets just 4 charges (weeks 12/24/36/48) → 12-week net LTV \$135.53, **blended \$119.74**. Since only 20% of the mix moves, blended net LTV sits in a tight **\$120–\$125 band** either way. (*Which horizon the reviewer intends is the one clarifying question worth asking.*)

Read-through: the **4-week plan carries the business** (70% × \$123.70); the **12-week plan is the most valuable per user** (\$162) but under-represented — a natural mix-shift target, which is exactly what Task 2 tests.

D.2 — A/B plan-upgrade break-even

Setup. Test group: the 4-week buyer is shown a **Plan Upgrade to 12-week** (\$49.99 one-time) instead of the second upsell; a buyer is billed under 12-week economics thereafter.

Intro and first upsell are identical in both arms (they cancel), and the 12% fee is uniform (cancels too), so break-even is clean:

$$p^* = E[\text{2nd upsell replaced}] / (\text{UpgradePrice} + \text{RecRev}_{12} - \text{RecRev}_4) \\ = 3.60 / (49.99 + 119.67 - 95.98) = 3.60 / 73.68 = 4.88\%$$

Each upgrade **adds ~\$73.68** (the \$49.99 surcharge + the richer 12-week recurring stream) and **cannibalises only \$3.60** (the expected second upsell).

UPGRADE TAKE-RATE	TEST NET LTV / 4-WK BUYER	LIFT VS CONTROL (\$123.70)
0% (control)	\$123.70	—
5%	\$123.77	+0.1%
10%	\$127.02	+2.7%
20%	\$133.50	+7.9%
30%	\$139.98	+13.2%

Break-even ≈ 4.9% (primary) to ≈ 8.3% under the conservative 52-week horizon.

Recommendation: run the test. The downside is tiny and bounded (the \$3.60 second upsell, test arm only); the upside is a permanent shift toward the highest-LTV plan. But because break-even is so low, the spreadsheet isn't the real question — the **live test must watch second-order effects the point-model ignores**: (a) does inserting an upgrade step depress overall checkout completion or first-upsell take? (b) do upgraded users **churn/refund faster** because they were pushed into a bigger

commitment? Track early-period retention and refunds on upgraded cohorts, not just the Day-0 take-rate.

The economics thesis (why retention is the moat)

This bridges Parts A, C, and D: it takes the per-plan LTV (Part D) up to the whole business, then shows *why* the Challenge (Part C) is an economic instrument, not just a UX feature.

Funnel reconciliation — 1,000 signups/day → annual net revenue

Built bottom-up, every stage tagged GIVEN / BENCHMARK / ASSUMPTION with a low/base/high band:

```
1,000 signups/day
  × ~60% quiz completion          (assumption)
  × ~92% complete → paywall view  (assumption)
  × ~10.7% paywall → purchase     (benchmark: hard-paywall trial→paid)
  ≈ 59 buyers/day
  × $118.75 net LTV/buyer          (blended $125 - 5% education-refund)
  × 365
  ≈ $2.56M / year (base)         range ≈ $1.19M - $4.52M
```

The two pins carrying the model are the **10.7% hard-paywall conversion** and the **AI-churn premium** (below). I'd plan on the **low-to-stress band (~\$1.2–2.1M)**, not the headline — our \$125 LTV is optimistic vs. RevenueCat's ~\$26 annual-realized median (ours is full-curve lifetime, but still rich).

The tail is the business — and AI apps churn through it faster

Decomposing the \$125 blended LTV: about **\$23 is intro cash**, which is retention-independent, and the remaining **84% arrives later** — 68% recurring plus 16% upsells. Only the intro is safe. That tail is exactly what the "AI apps churn ~30% faster" benchmark attacks:

- Applying the premium (monthly churn ~9.4% → ~12.2%, tail ×0.77) cuts blended LTV to **~\$101 (–19%)** and annual net revenue from \$2.56M to **~\$2.06M (–\$0.5M/yr)**.
- If the \$125 was itself built on a non-AI curve, **~\$2.06M is the truer base**.

Only about **16% of lifetime value is Day-1 cash**. The rest is hostage to a retention curve that AI apps are measurably worse at holding. That single fact is why the economics are *front-loaded and fragile* — and why the Challenge exists.

Retention beats the funnel (and it's safer)

LEVER	ANNUAL VALUE	NOTES
+1–2pt monthly retention	+\$160–350K/yr	Compounds across all cohorts; exceeds the entire refund drag; we start from the <i>weak</i> AI baseline (most headroom).
+1pt payroll conversion	+\$239K/yr	Near the ceiling already; raises refund / chargeback / ARL risk.
Cut refunds (edu ~5.1%)	~\$27K per point (~\$135K/yr base)	Directly recovered by better expectation-setting.
Stay web-billed vs app-store	+\$0.2–0.4M/yr	Distribution-fee avoidance.

Defend the tail; don't squeeze the funnel. The Challenge (Part C) is the instrument that buys back the AI-churn premium — a 1–2pt retention lift is worth more than a near-impossible conversion gain, and carries none of the regulatory downside.

The asymmetric tail risk — and the one move that fixes three problems

The regulatory exposure is not theoretical: **the FTC sued Genesis (a peer quiz-funnel studio) in June 2026**, and the reference **\$7.5M ARL settlement $\approx$ 2.9 $\times$ annual net revenue** — an existential, asymmetric tail. The drags also *stack* (refunds + app-store + churn).

Transparency + easy-cancel is a single move that fixes three things at once: it cuts refunds, defuses the regulatory existential risk, *and* lifts voluntary retention. It's the rare compliance action that is also the highest-ROI growth action.

Net thesis: the economics are front-loaded and fragile — a thin slice of certain Day-1 cash sitting on a fragile, fast-decaying recurring tail. So retention and trust are the moat, and the two initiatives that build them — the Challenge (habit) and transparency/easy-cancel (trust) — are not cost centers; they are the highest-leverage economic bets Jobscape can make.

Appendix

A — Jobscape quiz funnel, full decode (quiz_version v7.0.8)

Reconstructed from `api.funnel.jobscape.me/constructor/quiz-node/` — 37 pages, one branch (`used_ai` Yes/No). Questions and the variable each writes:

#	SECTION	QUESTION	OPTIONS	WRITES
0	—	Have you ever used Claude?	Yes / No $\rightarrow$ branch	<code>used_ai</code>
2	Profile	I want to learn Claude for...	Work / Personal / Growth	<code>why_claude</code>

#	SECTION	QUESTION	OPTIONS	WRITES
3	Profile	Current work status?	Full-time / Freelancer / Business owner / Between jobs / Exploring	status
4	Profile	How old are you?	18-24 / 25-34 / 35-44 / 45-54 / 55+	age
5	Profile	Gender identity?	Female / Male / Skip	gender
6	Profile	How would learning Claude benefit you?	Promotion / Work faster / Confidence / Start business / Earn more	goal
8	Challenges	Rate your experience with AI so far	Great / Good-but / Frustrating / Haven't tried	experience_ai
9	Challenges	Ready for how AI is changing careers?	Need skills fast / Bit worried / Somewhat / Confident	skills_competitive
10	Challenges	What scares you most about AI and your career?	Replaced by AI-savvier peer / Falling behind / Losing opportunities / Nothing	career_scares
12	Challenges	I create documents/reports regularly	1-5 scale	often_create
13	Challenges	Used AI to write, then spent as long fixing it?	Every time / Sometimes / No / Haven't	long_fix
14	Challenges	What would change with polished docs in 2 min?	3x content / stop procrastinating / look pro / strategy	polished_documents
16	Challenges	Had an app idea but couldn't build it?	Yes / No	idea_app
17	Challenges	What stopped you?	Can't code / dev too expensive / gave up / didn't know start	app_stoppers
18	Challenges	Possible to build an app without coding?	No / Maybe w/ AI / Yes	app_without
20	Personalization	How do you prefer to learn?	Own pace / Set deadlines	prefer_learn
21	Personalization	Time to dedicate to your goal?	10 / 20 / 30 min / 1 hr per day	time_goal
22	Personalization	Approach that works best?	80% theory / 80% practice	approach_best
23	Personalization	Include projects in a portfolio site?	Yes / No	coding_experience
24	Personalization	Want an AI mentor to guide you?	Yes / No	ai_mentor_prefer
26	Personalization	Would an AI certification help your career?	Definitely / Probably / No	certification_advantage

#	SECTION	QUESTION	OPTIONS	WRITES
28	Personalization	What kept you from advancing before?	No plan / fear waste time / too busy / too confusing	stopped_before
31	—	Email capture	—	email
32	—	Name capture	—	name
34– 35	Sell	Selling page + "chase" offer	—	paywall

Interleaved teasers (the persuasion layer): FOMO cards after the `used_ai` branch; *"There is nothing to worry about"* (+ Harvard "humans with AI will replace humans without AI") after `career_scares`; *"Create polished content"* after the document pain; *"Building an app has never been this easy"* after the app-build pain; a certificate teaser; an "AI profile" reveal; and a fake-progress loader before the paywall. Full raw map: `part-a-audience/materials/quiz-map.md`.

B — Reproducibility

- **Quiz decode:** open funnel constructor API (no auth); 37 nodes pulled and parsed → `quiz-map.md` + `quiz-nodes-raw.json`.
- **Economics:** `part-d-economics/model/ltv_model.py` reads `data/plans.csv`, prints both horizon readings and the A/B scenarios.
- **Research:** 10 sourced agent reports (5 broad + 5 deep, incl. product-reality from public reviews) under `research/`; every external claim carries a URL in the underlying files.